

Approximate CNN

M.Abdelmaseeh, M.Romani, K.Khallaf, M.Tarek and A.Ibrahim

December 2022

Github Link : <https://github.com/MinaAbdelmaseeh/Approximate-CNN>


1 Introduction

Approximate computing is a powerful way to decrease the consumption power of digital integrated circuits while maintaining a decent accuracy. In this project a Convolutional neural network is implemented using both accurate and approximate paradigms. Additionally, the accuracy, the power and the area are compared.

2 Accurate 2x2 Multiplier

The implementation of the accurate multiplier providing the multiplication output of two numbers, where each number consists of two bits, is constructed using the presented configuration in the research paper [1] (such that the multiplication process is accomplished using 6 AND gates and 2 XOR gates). Then, the constructed 2bit*2bit multiplier is used to represent the building block for implementing an accurate 4bit*4bit multiplier and an accurate 8bit*8bit multiplier.

Accurate 2x2 Multiplier testbench results:



3 Approximate 2x2 Multiplier

the approximate 2x2 multiplier is implemented by reducing the output size from 4 bits to 3 bits so the result of

$$(11)_2 * (11)_2$$

is

$$(111)_2$$

instead of

$$(1001)_2$$

. By this, we can save some power by reducing the number of logic gates used, but in the same time we sacrifice some precision and accuracy. the multiplier uses only 4 and gates and one or gate.

Approximate 2x2 Multiplier testbench results:




4 Building larger Multipliers

4.1 Accurate Multipliers

4.1.1 4x4 Accurate Multiplier

The accurate 4bit*4bit multiplier is constructed using the previously implemented 2bit*2bit accurate multiplier such that the output of four partial products represent the result of the multiplication process. Then, the sum of these four partial products is accomplished using a Wallace tree where the last stage is implemented using a carry ripple adder. The Wallace tree is implemented using two full-adders and two half-adders representing the first stage and another two full adders and two half adders for the second stage, while the last stage is implemented using a 4bit carry ripple adder.


4x4 Accurate Multiplier testbench results:



4.1.2 8x8 Accurate Multiplier

The accurate 8bit*8bit multiplier is implemented using the illustrated 4bit*4bit accurate multiplier to produce the required partial products. Similarly, the sum of the partial products is accomplished using a Wallace tree, where the first stage is implemented using four full adders and four half adders, the second stage is implemented using another four full adders and four half adders, and the last stage is represented by a 10bit carry ripple adder.

8x8 Accurate Multiplier testbench results:




4.2.1 4x4 Approximate Multiplier

The diagram illustrates the carry propagation in a 3-bit ripple-carry adder. It shows three rows of dots representing bits. The top row has 6 dots, the middle row has 6 dots, and the bottom row has 10 dots. The first three dots of the bottom row are grouped by a box and labeled '3-bit adder'. Arrows show the carry propagating from the first carry-in to the second, third, and fourth carry-in positions.

3


4x4 Approximate Multiplier testbench results:



4.2.2 8x8 Approximate Multiplier

the 8x8 multiplier is implemented by using the 4x4 approximate multipliers as sub-units but the Wallace tree implementation is the same as the accurate 8x8 multiplier as the output of the 4x4 approximate multiplier is 8 bits which is the same size of the output of the accurate one. The first stage is implemented using four full adders and four half adders, the second stage is implemented using another four full adders and four half adders, and the last stage is represented by a 10-bit carry ripple adder.

8x8 Accurate Multiplier testbench results:



5 Fixed Point Representation

In our implementation, an 8-bit fixed point numbers are used. we assigned 4 bits for the integer part and 4 bits for the decimal part. On multiplying two 8-bit fixed point numbers the result is as follows:


$$X_{(7:4)} \cdot X_{(3:0)} \times Y_{(7:4)} \cdot Y_{(3:0)} = M_{(15:8)} \cdot M_{(7:0)} \quad (1)$$

On multiplying two 8-bit numbers, the result is represented in 16 bits. However, in fixed point representation the result can be truncated with a loss in precision. The result is then represented as $M_{(11:8)} \cdot M_{(7:4)}$.

In order to take into account the overflow case, before the truncation process of the result from 16 bit to 8 bit, the number is checked if it was greater than the maximum representable number in 8 bits and the result is then the maximum representable number(either positive or negative).

The maximum representable positive number is "0111.1111", while the maximum representable negative number is "1000.0000".

6 Convolutional Neural Network



The convolutional neural network is composed of multiple neural layers, each has an set of inputs (n) equal to the number of neurons in that layer and a set of outputs (m) equal to the inputs of the next layer. The neural layer can then be represented as the following equation

$$f\left(\begin{bmatrix} W_{00} & \dots & W_{0n} \\ \vdots & \ddots & \vdots \\ W_{0m} & \dots & W_{mn} \end{bmatrix} \cdot \begin{bmatrix} In_0 \\ \vdots \\ In_n \end{bmatrix} + \begin{bmatrix} B_0 \\ \vdots \\ B_n \end{bmatrix}\right) = \begin{bmatrix} Out_0 \\ \vdots \\ Out_m \end{bmatrix} \quad (2)$$

where $[W]$ is the weight matrix, $[B]$ is the bias vector, $[In]$ is the input to the neural layer, $[Out]$ is the output of the neural layer and f is the nonlinear activation function.

6.1 Activation function

for our implementation we used basic ReLU function defined as follows :

$$f(x) = \begin{cases} 0 & \text{if } x \leq 0 \\ x & \text{if } x > 0 \end{cases} \quad (3)$$

This can be easily implemented in 8-fixed point representation by checking on the most significant bit if it is equal to 1 "negative number" output 0 ,otherwise output = input.

7 Results

7.1 Functional testing : Test Case

for a neural network of 4 layers (2 hidden layers) : 2 - 4 - 4 - 2 :

$$In = \begin{bmatrix} 0.C \\ 0.9 \end{bmatrix}_{16} = \begin{bmatrix} 0.75 \\ 0.5625 \end{bmatrix}_{10} \quad (4)$$

Layer 2 :

$$W_2 = \begin{bmatrix} 0.7 & 0.7 \\ 0.7 & 0.7 \\ 0.7 & 0.7 \\ 0.7 & 0.7 \end{bmatrix}_{16}, B_2 = \begin{bmatrix} 0.7 \\ 0.7 \\ 0.7 \\ 0.7 \end{bmatrix}_{16} \quad (5)$$

Layer 3 :

$$W_3 = \begin{bmatrix} 0.7 & 0.7 & 0.7 & 0.7 \\ 0.8 & 0.8 & 0.8 & 0.8 \\ 0.8 & 0.8 & 0.8 & 0.8 \\ 0.8 & 0.8 & 0.8 & 0.8 \end{bmatrix}_{16}, B_3 = \begin{bmatrix} 0.8 \\ 0.8 \\ 0.8 \\ 0.8 \end{bmatrix}_{16} \quad (6)$$

Layer 4 :

$$W_4 = \begin{bmatrix} 0.8 & 0.8 & 0.8 & 0.8 \\ 0.8 & 0.8 & 0.8 & 0.8 \end{bmatrix}_{16}, B_4 = \begin{bmatrix} 0.8 \\ 0.8 \end{bmatrix}_{16} \quad (7)$$

$$\text{Expected Out} = \begin{bmatrix} 5.6 \\ 5.6 \end{bmatrix}_{16} = \begin{bmatrix} 5.4204 \\ 5.4204 \end{bmatrix}_{10} \quad (8)$$

$$\text{Simulated Out (Accurate Neural network)} = \begin{bmatrix} 4.E \\ 4.E \end{bmatrix}_{16} = \begin{bmatrix} 4.875 \\ 4.875 \end{bmatrix}_{10} \quad (9)$$

$$\text{Simulated Out (Approx. Neural network)} = \begin{bmatrix} 4.C \\ 4.C \end{bmatrix}_{16} = \begin{bmatrix} 4.75 \\ 4.75 \end{bmatrix}_{10} \quad (10)$$

Note that the expected output is slightly different than the simulated out, this is due to the truncation of the 8x8 multiplier which results in a decrease in accuracy.

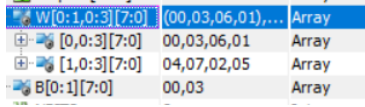
7.2 Functional Testing : more test cases ...

| | | |
|-----------------|-------------------|-------|
| W[0:3,0:1][7:0] | (08,05),(02,07... | Array |
| [0,0:1][7:0] | 08,05 | Array |
| [1,0:1][7:0] | 02,07 | Array |
| [2,0:1][7:0] | 04,01 | Array |
| [3,0:1][7:0] | 06,03 | Array |
| B[0:3][7:0] | 00,01,02,03 | Array |

Figure 2: The Weights and Biases of the 1st layer

| | | |
|-----------------|-------------------|-------|
| W[0:3,0:3][7:0] | (04,07,02,05),... | Array |
| [0,0:3][7:0] | 04,07,02,05 | Array |
| [1,0:3][7:0] | 00,03,06,01 | Array |
| [2,0:3][7:0] | 04,07,02,05 | Array |
| [3,0:3][7:0] | 00,03,06,01 | Array |
| B[0:3][7:0] | 04,07,02,05 | Array |

Figure 3: The Weights and Biases of the 2nd layer



| | | |
|-----------------|-------------------|-------|
| W[0:1,0:3][7:0] | (00,03,06,01),... | Array |
| [0,0:3][7:0] | 00,03,06,01 | Array |
| [1,0:3][7:0] | 04,07,02,05 | Array |
| B[0:1][7:0] | 00,03 | Array |

Figure 4: The Weights and Biases of the 3rd layer



| | | | | | | | |
|-----------------------|------|------|------|------|------|------|------|
| Inputs[15:0] | 0502 | 0e09 | 0502 | 0707 | 0b0d | fb0d | fbfd |
| output_accurate[15:0] | 0802 | 0b04 | 0802 | 0802 | 0b04 | 0802 | 0801 |
| output_approx[15:0] | 0702 | 0a03 | 0702 | 0802 | 0b04 | 0802 | 0701 |
| write[7:0] | 03 | | | | 03 | | |

Figure 5: The result of the simulation. As shown in the figure, some cases occur where the output of the approximate neural network is different than the output of the accurate neural network.

7.3 Power, Area and frequency Analysis

For 8-bit multipliers, In the following table a comparison between the accurate and the approximate 8x8 multiplier is presented, As shown in the table, the approximate multiplier consume less Dynamic power by 7% and less LUTs by 46%. The Frequency chosen to operate the multipliers is 100 MHz in order to ensure low power consumption :

| - | Accurate | Approximate |
|---------------|----------|-------------|
| Static power | 0.081 | 0.082 W |
| Dynamic power | 0.076 | 0.071 W |
| Total power | 0.157 W | 0.152 W |
| LUTs | 114 | 78 |
| Frequency | 100 MHz | 100 MHz |

For 2 - 4 - 4 - 2 neural network, The approximate power consumption is less by 6% and less LUTs by 10% than the accurate neural network and less area by 10%. The Frequency chosen to operate the neural network is 100 MHz in order to ensure low power consumption:

| - | Accurate | Approximate |
|---------------|----------|-------------|
| Static power | 0.081 W | 0.081 W |
| Dynamic power | 0.098 W | 0.088 W |
| Total power | 0.179 W | 0.169 W |
| LUTs | 3347 | 3037 |
| Registers | 336 | 336 |
| Frequency | 100 MHz | 100 MHz |

For 2 - 10 - 10 - 2 neural network, The approximate power consumption is less by 7% and less LUTs by 13% than the accurate neural network and less area by 9%. The Frequency chosen to operate the neural network is 100 MHz in order to ensure low power efficiency:

| - | Accurate | Approximate |
|---------------|----------|-------------|
| Static power | 0.082 W | 0.082 W |
| Dynamic power | 0.392 W | 0.357 W |
| Total power | 0.474 W | 0.439 W |
| LUTs | 14206 | 12568 |
| Registers | 1296 | 1296 |
| Frequency | 100 MHz | 100 MHz |

7.4 Accuracy estimation

In order to estimate the accuracy of approximate neural network, we initiated several pseudo random neural networks and gathered the following results:

| Input | Accurate output | Approximate output |
|-------------|-----------------|--------------------|
| 0x0.c 0x0.9 | 0x0.7 0x0.9 | 0x0.6 0x0.9 |
| 0x0.5 0x0.2 | 0x0.6 0x0.7 | 0x0.6 0x0.7 |
| 0x0.7 0x0.7 | 0x0.6 0x0.7 | 0x0.7 0x0.9 |
| 0x0.b 0x0.d | 0x0.7 0x0.9 | 0x0.6 0x0.9 |
| 0xf.b 0x0.d | 0x0.6 0x0.6 | 0x0.6 0x0.6 |
| 0xf.b 0xf.d | 0x0.6 0x0.6 | 0x0.6 0x0.6 |
| 0x2.3 0x1.6 | 0x0.8 0x0.d | 0x0.8 0x0.c |

The error is calculated as:

$$Err = \frac{|\mathbf{accurate} - \mathbf{approx}|}{\sum \mathbf{accurate}} \quad (11)$$

From the above table , the error is estimated to be around 5.82%.

Note : The results are checked using a python code that simulates the accurate neural network. The code can be found in the Github link in testbenches folder.

8 Conclusion

As it was shown in the results section, the power efficiency difference between the approximate neural networks and the accurate neural network increases as the system gets more complex. Hence, Approximate neural network proved to decrease the power consumption and area utilization by at least 10% for small circuits.

Although the ratio between the utilization of approximate and the accurate implementations of 8x8 multipliers is around 50%, The power consumption seemed to stall at 10-13% for small circuits. This is because FPGAs uses LUTs so the reduction in the schematic might not be very powerful when implemented using FPGAs. However, It is expected to have better power enhancement when implemented as an ASIC.

9 References

- [1] P. Kulkarni, P. Gupta, and M. Ercegovac. 2011. Trading Accuracy for Power in a Multiplier Architecture. *J.Low Power Electronics* 7, 4, 490-501.